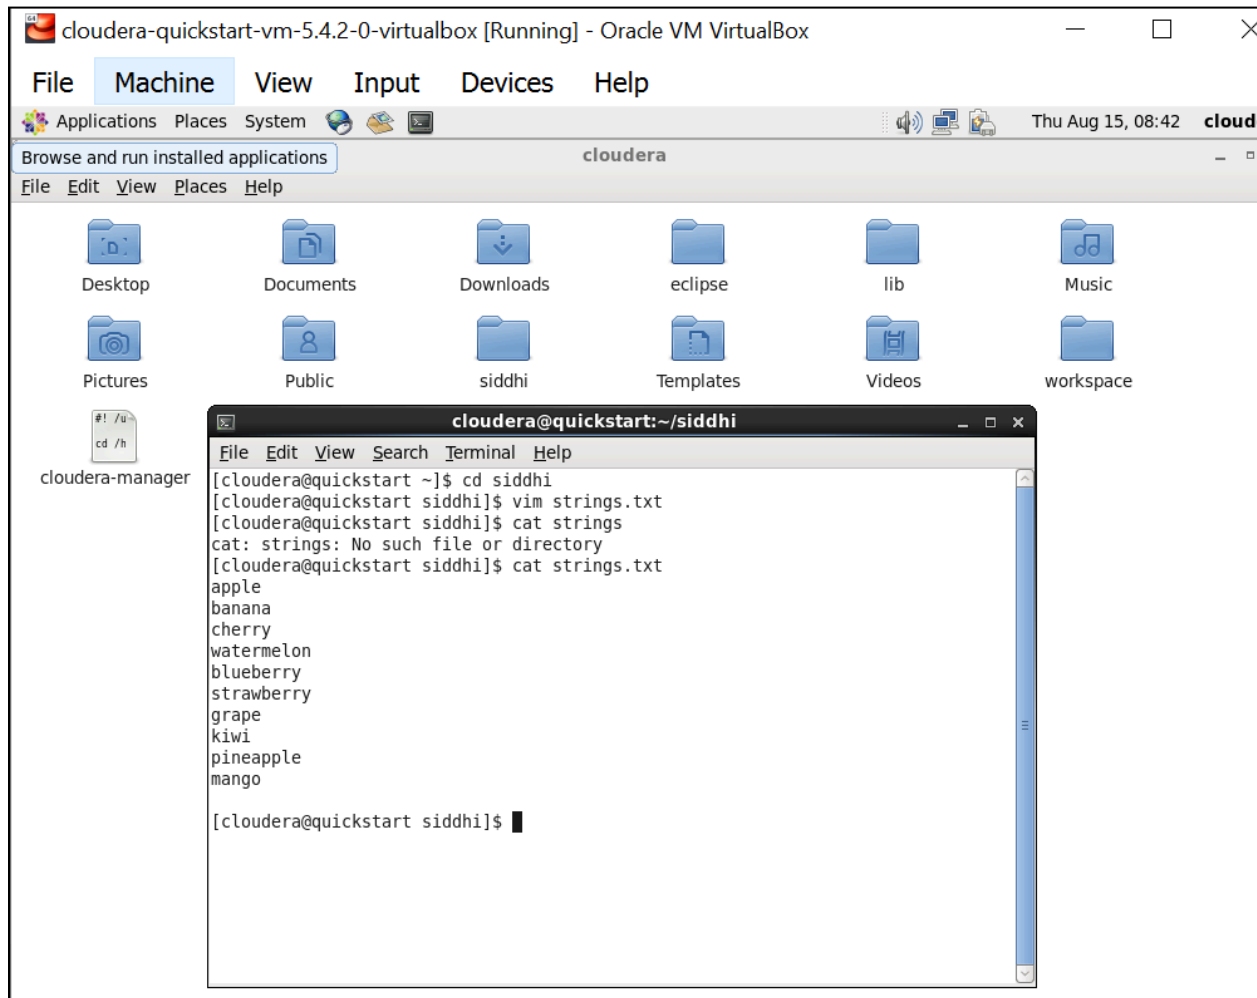


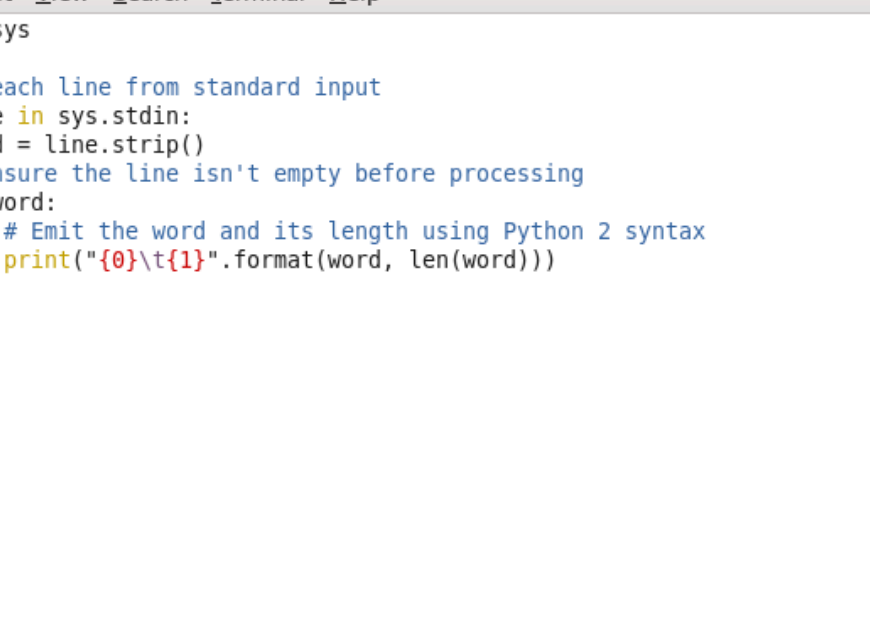
Big Data Analytics and Visualization

Lab 2 **MapReduce**

Aim: 1. Write a mapper and reducer to find the longest string in a text file. (strings.txt, create one with at least 10 words)

Output:





```
cloudera@quickstart:~/siddhi
File Edit View Search Terminal Help
import sys

# Read each line from standard input
for line in sys.stdin:
    word = line.strip()
    # Ensure the line isn't empty before processing
    if word:
        # Emit the word and its length using Python 2 syntax
        print("{0}\t{1}".format(word, len(word)))

"mapper.py" 11L, 274C 1,1 All
```

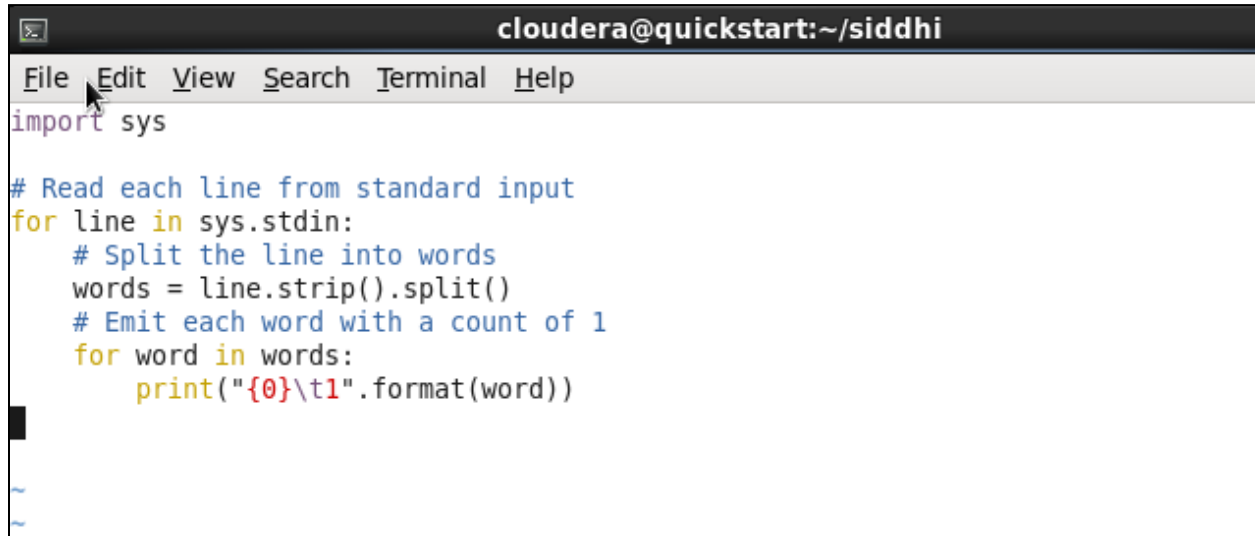
```
[cloudera@quickstart siddhi]$ vim reducer.py
[cloudera@quickstart siddhi]$ vim mapper.py
[cloudera@quickstart siddhi]$ cat mapper_output.txt | python reducer.py
The longest word is 'watermelon' with length 10
[cloudera@quickstart siddhi]$ █
```

Aim: 2. Write the following Map Reduce program to understand Map Reduce Paradigm. (Create your own .txt file)

- a. WordCount
- b. Average WordCount
- c. Word with minimum count and word with maximum count

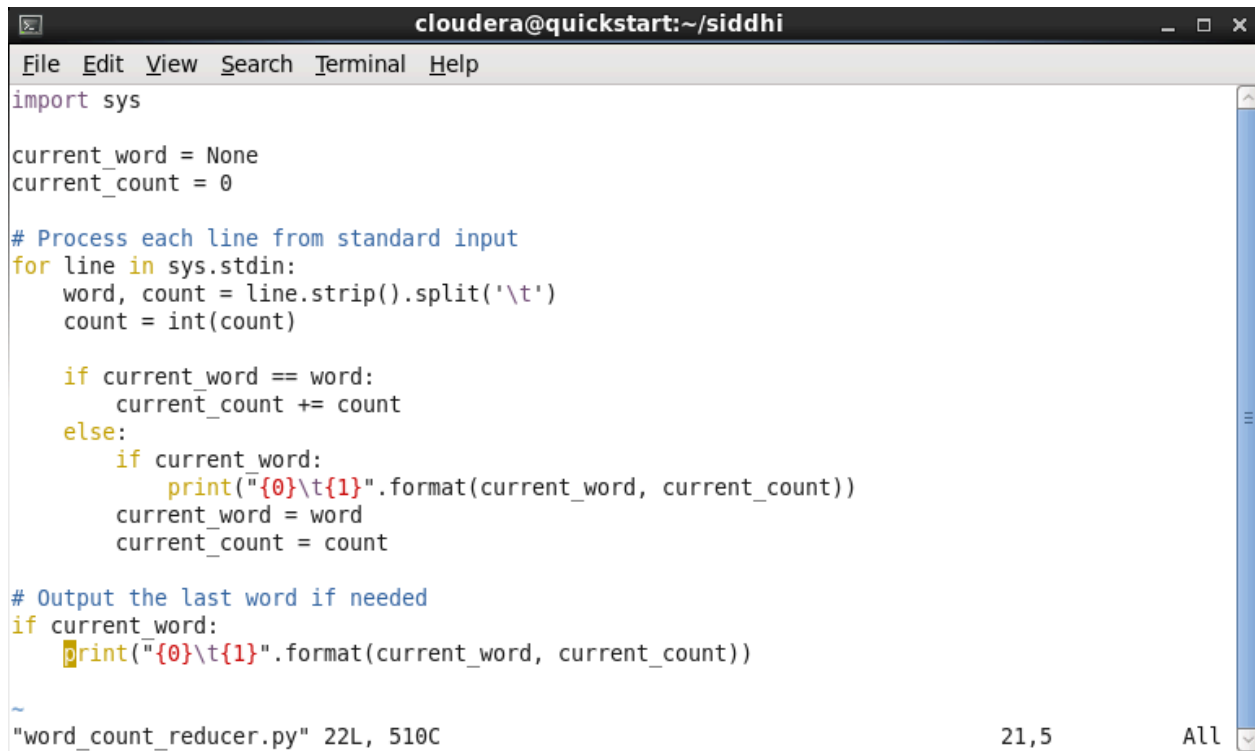
Output:

a. WordCount



```
cloudera@quickstart:~/siddhi
File Edit View Search Terminal Help
import sys

# Read each line from standard input
for line in sys.stdin:
    # Split the line into words
    words = line.strip().split()
    # Emit each word with a count of 1
    for word in words:
        print("{0}\t1".format(word))
```



```
cloudera@quickstart:~/siddhi
File Edit View Search Terminal Help
import sys

current_word = None
current_count = 0

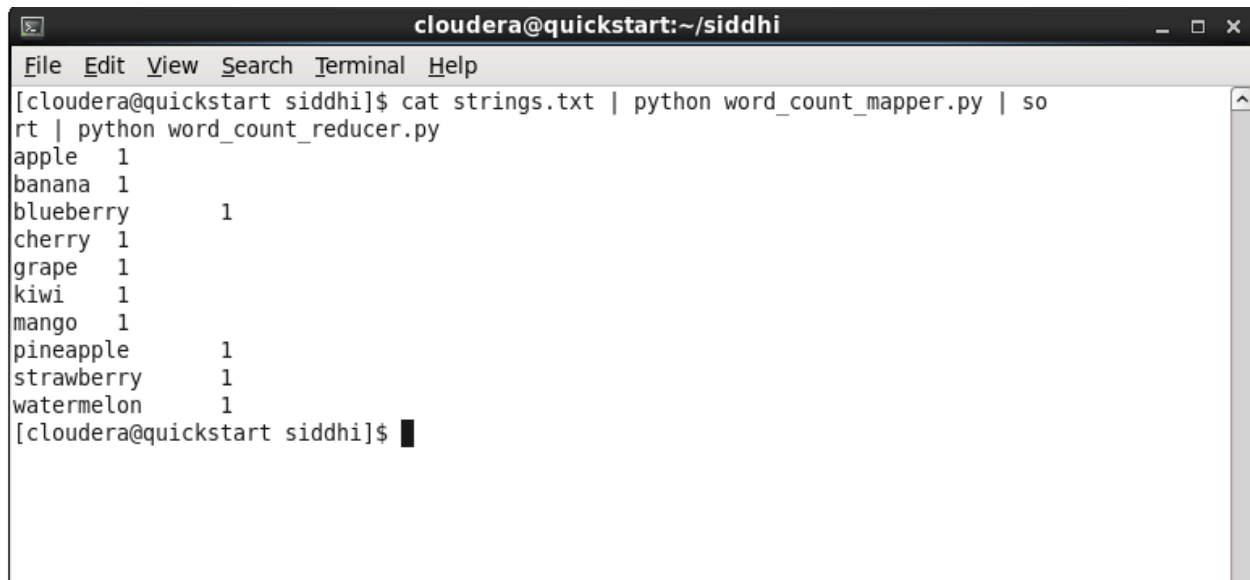
# Process each line from standard input
for line in sys.stdin:
    word, count = line.strip().split('\t')
    count = int(count)

    if current_word == word:
        current_count += count
    else:
        if current_word:
            print("{0}\t{1}".format(current_word, current_count))
            current_word = word
            current_count = count

# Output the last word if needed
if current_word:
    print("{0}\t{1}".format(current_word, current_count))

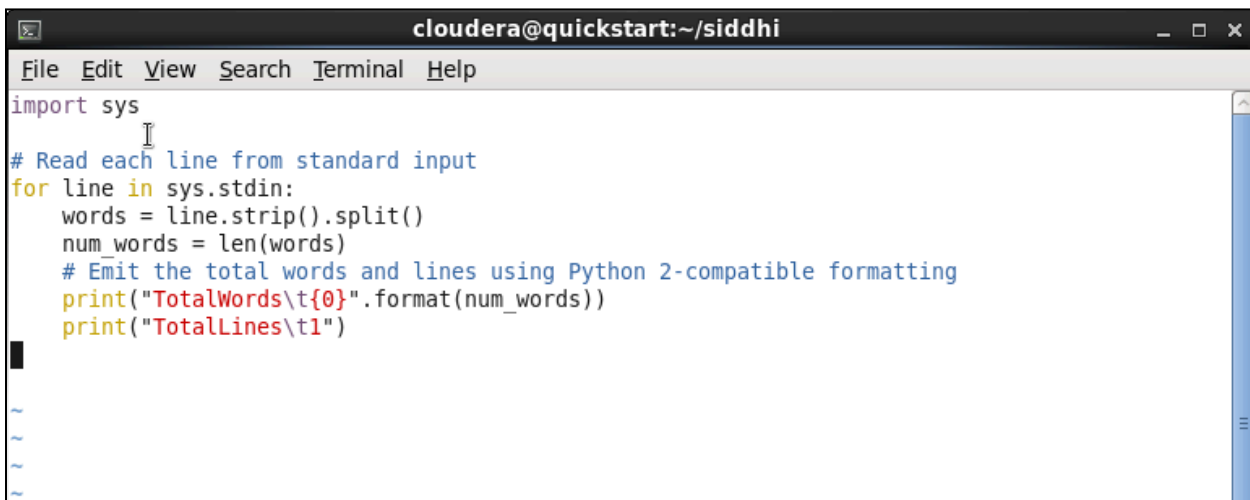
"word_count_reducer.py" 22L, 510C 21,5 All
```

Output:

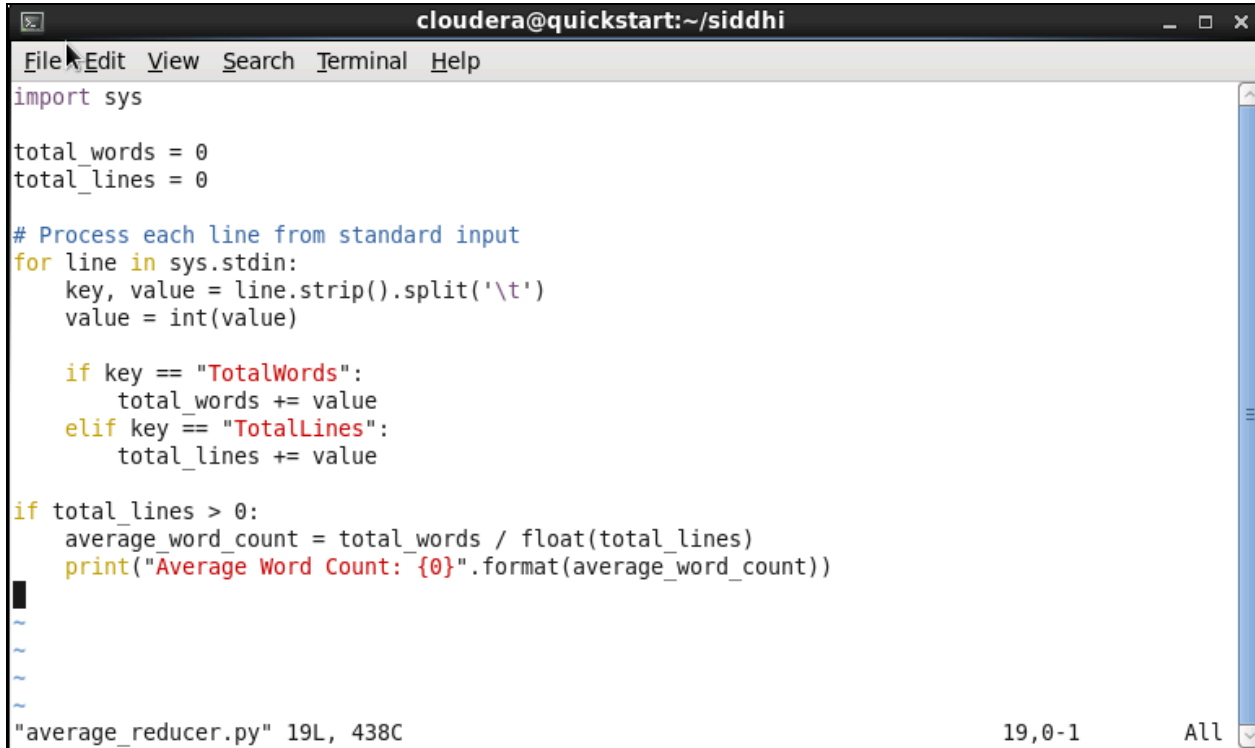


```
cloudera@quickstart:~/siddhi
File Edit View Search Terminal Help
[cloudera@quickstart siddhi]$ cat strings.txt | python word_count_mapper.py | so
rt | python word_count_reducer.py
apple 1
banana 1
blueberry 1
cherry 1
grape 1
kiwi 1
mango 1
pineapple 1
strawberry 1
watermelon 1
[cloudera@quickstart siddhi]$
```

b. Average WordCount



```
cloudera@quickstart:~/siddhi
File Edit View Search Terminal Help
import sys
# Read each line from standard input
for line in sys.stdin:
    words = line.strip().split()
    num_words = len(words)
    # Emit the total words and lines using Python 2-compatible formatting
    print("TotalWords\t{0}".format(num_words))
    print("TotalLines\t1")
```



The screenshot shows a text editor window titled "cloudera@quickstart:~/siddhi". The menu bar includes File, Edit, View, Search, Terminal, and Help. The code is as follows:

```
import sys

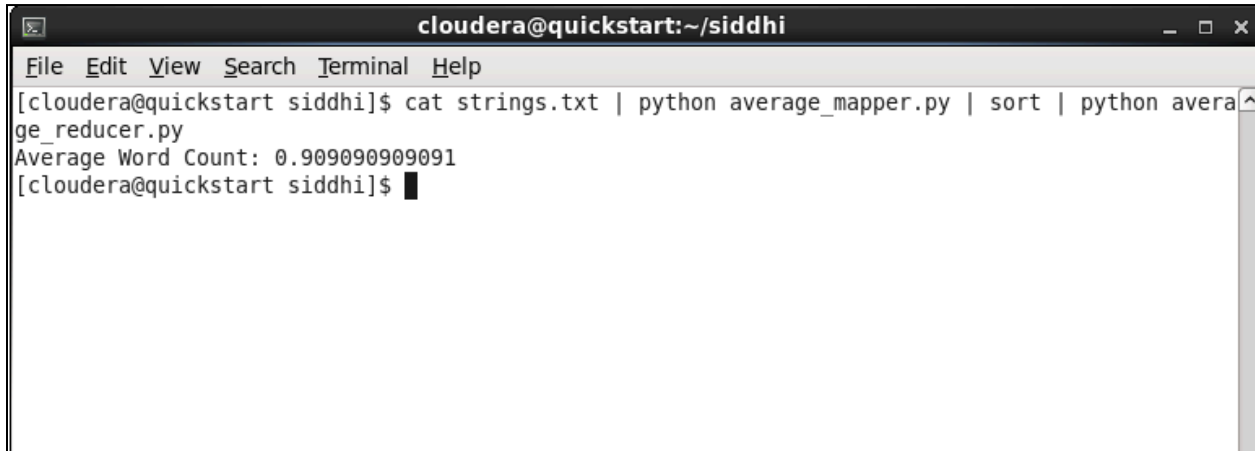
total_words = 0
total_lines = 0

# Process each line from standard input
for line in sys.stdin:
    key, value = line.strip().split('\t')
    value = int(value)

    if key == "TotalWords":
        total_words += value
    elif key == "TotalLines":
        total_lines += value

if total_lines > 0:
    average_word_count = total_words / float(total_lines)
    print("Average Word Count: {0}".format(average_word_count))
```

The status bar at the bottom indicates the file is "average_reducer.py" with 19 lines and 438 characters. The cursor is at line 19, column 1.



The screenshot shows a terminal window titled "cloudera@quickstart:~/siddhi". The command prompt is "[cloudera@quickstart siddhi]\$". The command entered is "cat strings.txt | python average_mapper.py | sort | python average_reducer.py". The output is "Average Word Count: 0.909090909091". The prompt is now "[cloudera@quickstart siddhi]\$".

c. Word with minimum count and word with maximum count



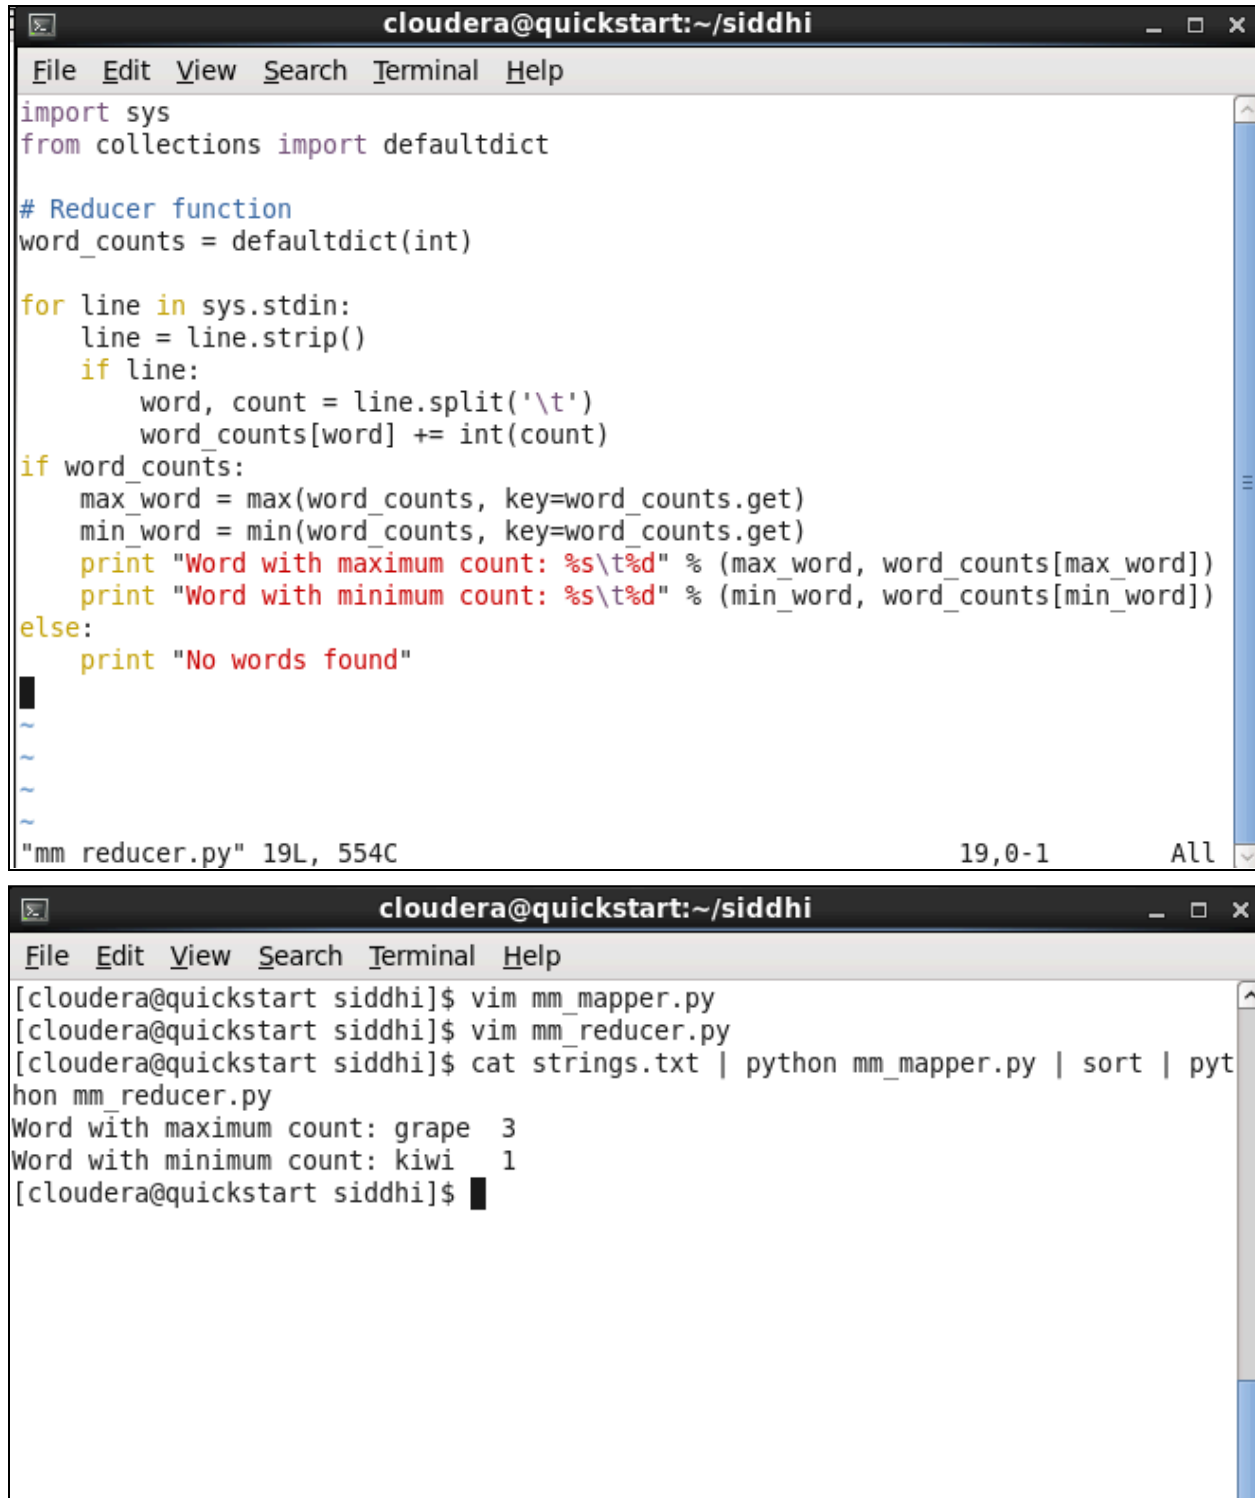
The image shows a terminal window titled "cloudera@quickstart:~/siddhi". The window contains a Python script for a Mapper function. The script imports the sys module and iterates over lines of input. Each line is stripped and split into words. For each word, it prints the word and its count (1). The terminal shows the script being executed on a file named "mm_mapper.py", which has 11 lines and 184 characters. The output shows the first line of the file being processed, with the word "word" appearing 1 time.

```
cloudera@quickstart:~/siddhi
File Edit View Search Terminal Help
#!/usr/bin/env python

import sys

# Mapper function
for line in sys.stdin:
    line = line.strip()
    words = line.split()
    for word in words:
        print "%s\t%d" % (word, 1)

"mm_mapper.py" 11L, 184C 11,0-1 All
```



The image consists of two screenshots of a terminal window titled "cloudera@quickstart:~/siddhi".

The top screenshot shows the code for a Python reducer program, "mm_reducer.py". The code imports the sys module and defaultdict from the collections module. It defines a reducer function that reads input lines, splits them by tab, and updates a defaultdict with word counts. It then finds the word with the maximum and minimum counts and prints them. The status bar at the bottom indicates "mm_reducer.py" is 19 lines long and 554 characters.

```
import sys
from collections import defaultdict

# Reducer function
word_counts = defaultdict(int)

for line in sys.stdin:
    line = line.strip()
    if line:
        word, count = line.split('\t')
        word_counts[word] += int(count)
if word_counts:
    max_word = max(word_counts, key=word_counts.get)
    min_word = min(word_counts, key=word_counts.get)
    print "Word with maximum count: %s\t%d" % (max_word, word_counts[max_word])
    print "Word with minimum count: %s\t%d" % (min_word, word_counts[min_word])
else:
    print "No words found"
```

The bottom screenshot shows the terminal commands used to edit the files, view the input data, and run the program. The output shows the word "grape" with a count of 3 and the word "kiwi" with a count of 1.

```
[cloudera@quickstart siddhi]$ vim mm_mapper.py
[cloudera@quickstart siddhi]$ vim mm_reducer.py
[cloudera@quickstart siddhi]$ cat strings.txt | python mm_mapper.py | sort | python mm_reducer.py
Word with maximum count: grape 3
Word with minimum count: kiwi 1
[cloudera@quickstart siddhi]$
```


Aim: 3. Implement matrix multiplication with Hadoop Map Reduce

Output:

```
cloudera@quickstart:~/workspace/siddhi
File Edit View Search Terminal Help
[cloudera@quickstart workspace]$ cd siddhi
[cloudera@quickstart siddhi]$ cat mat.txt
A,0,0,1
A,0,1,2
A,0,2,3
A,1,0,4
A,1,1,5
A,1,2,6
A,2,0,7
A,2,1,8
A,2,2,9
B,0,0,9
B,0,1,8
B,0,2,7
B,1,0,6
B,1,1,5
B,1,2,4
B,2,0,3
B,2,1,2
B,2,2,1
```

```
cloudera@quickstart:~/workspace/siddhi
File Edit View Search Terminal Help
B,2,1,2
B,2,2,1
[cloudera@quickstart siddhi]$ cat mat.txt |python matmap.py
A      0,0,1
A      0,1,2
A      0,2,3
A      1,0,4
A      1,1,5
A      1,2,6
A      2,0,7
A      2,1,8
A      2,2,9
B      0,0,9
B      0,1,8
B      0,2,7
B      1,0,6
B      1,1,5
B      1,2,4
B      2,0,3
B      2,1,2
B      2,2,1
```

```
[cloudera@quickstart siddhi]$ cat mat.txt | python matmap.py | python matreducer.py
(0, 0) 30
(0, 1) 24
(0, 2) 18
(1, 0) 84
(1, 1) 69
(1, 2) 54
(2, 0) 138
(2, 1) 114
(2, 2) 90
[cloudera@quickstart siddhi]$
```

Aim: 4. Amazon collects item sold data every hour at many locations across the globe and gathers a large volume of log data, which is a good candidate for analysis with Map Reduce since it is semi-structured and record-oriented. Write a Map Reduce program that sorts unit sold data. (Refer ordered_unitsold_data.txt file)

Output:

```
cloudera@quickstart:~/workspace/siddhi
File Edit View Search Terminal Help
[cloudera@quickstart siddhi]$ cat amazonsales.txt
InvoiceNo,StockCode,Description,Quantity,InvoiceDate,UnitPrice,CustomerID,Country
10001,85123A,White Hanging Heart T-Light Holder,6,12/1/2024 8:26,2.55,17850,United Kingdom
10002,71053,White Metal Lantern,6,12/1/2024 8:28,3.39,17850,United Kingdom
10003,84406B,Cream Cupid Hearts Coat Hanger,8,12/1/2024 8:34,2.75,13047,France
10004,84029G,Knitted Union Flag Hot Water Bottle,6,12/1/2024 8:35,3.39,13047,France
10005,84029E,Red Woolly Hottie White Heart,6,12/1/2024 8:36,3.39,13047,France
10006,22752,Set of 3 Butterfly Cookie Cutters,12,12/1/2024 8:40,2.75,12583,United States
10007,21730,Glass Star Frosted T-Light Holder,6,12/1/2024 8:41,4.25,12583,United States
10008,71053,White Metal Lantern,6,12/1/2024 8:45,3.39,13748,Germany
10009,84406G,Cream Hearts Coat Hanger,8,12/1/2024 8:48,2.75,15100,Australia
10010,84029E,Red Woolly Hottie White Heart,6,12/1/2024 8:50,3.39,15100,Australia

[cloudera@quickstart siddhi]$ cat amazonsales.txt | python amazonmapper.py
United Kingdom 15
United Kingdom 20
France 22
France 20
France 20
United States 33
United States 25
Germany 20
Australia 22
Australia 20

[cloudera@quickstart siddhi]$ cat amazonsales.txt | python amazonmapper.py | python amazonreducer.py
United Kingdom 35.0
United States 58.0
Australia 42.0
Germany 20.0
France 62.0
[cloudera@quickstart siddhi]$
```